

DL_Lab1

September 18, 2020

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
%matplotlib inline

[288]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[ ]: plt.style.use(['ggplot'])

[ ]: df = pd.read_csv("/content/drive/My Drive/Deep Learning/Lab1/Admission_Predict.
→csv",engine='python')
df.columns.values

[ ]: array(['Serial No.', 'GRE Score', 'TOEFL Score', 'University Rating',
'SOP', 'LOR', 'CGPA', 'Research', 'Chance of Admit'], dtype=object)
```

1 Linear Regression With one variable

```
[ ]:

[ ]: df1 = pd.DataFrame(df, columns=[df.columns.values[1]]+[df.columns.values[8]])
df1

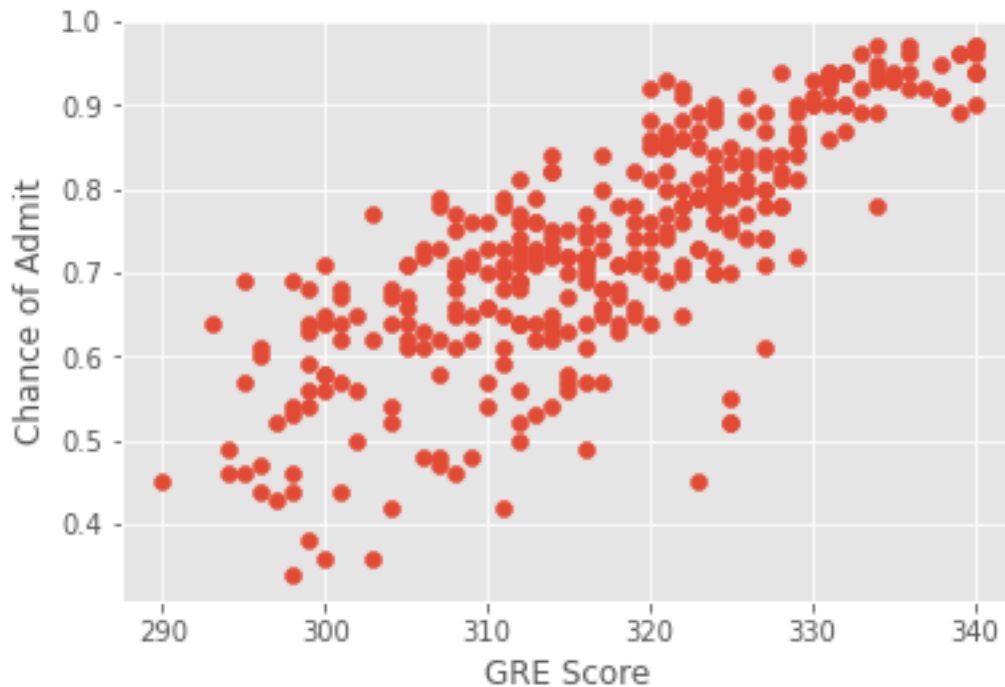
[ ]:
  GRE Score  Chance of Admit
0         337             0.92
1         324             0.76
2         316             0.72
3         322             0.80
4         314             0.65
..         ...             ...
395        324             0.82
396        325             0.84
397        330             0.91
398        312             0.67
```

399 333 0.95

[400 rows x 2 columns]

1.1 Plotting of Data:

```
[ ]: X = np.array([df1.values[:340,0]]).T
y = np.array([df1.values[:340,1]]).T
plt.scatter(X,y)
plt.xlabel(df1.columns.values[0])
plt.ylabel(df1.columns.values[1])
plt.show()
print(X.shape, y.shape)
```



(340, 1) (340, 1)

1.2 Gradient Descent:

```
[ ]: def cal_cost(theta,X,y):
    m = len(y)
    predictions = X.dot(theta)
    cost = (1/2*m) * np.sum(np.square(predictions-y))
    return cost
```

```
[ ]: def gradient_descent(X,y,theta,learning_rate=0.01,iterations=100):
    m = len(y)
    cost_history = np.zeros(iterations)
    theta_history = np.zeros((iterations,2))
    for it in range(iterations):
        prediction = np.dot(X,theta)
        # print(prediction.shape, y.shape)
        # print(np.array(prediction-y).shape)
        # print( X.T.dot((prediction - y)))
        theta = theta -(1/m)*learning_rate*( X.T.dot((prediction - y)))
        theta_history[it,:] =theta.T
        cost_history[it] = cal_cost(theta,X,y)
    return theta, cost_history, theta_history
```

```
[ ]: # theta = np.random.randn(2,1)
theta = np.array([[0],[0.1]])
learning_rate = 0.1
iterations = 1000
mean_X = np.sum(X)/len(X)
range_X = max(X)-min(X)
X1 = (X-mean_X)/range_X
X_b = np.c_[np.ones(len(X)),X1]
# print(theta.shape, X_b[0], y[0])
# print("sas",X_b.shape)
theta,cost_history,theta_history =_
    ↪gradient_descent(X_b,y,theta,learning_rate,iterations)

print('Theta0:          {:0.3f},\nTheta1:          {:0.3f}'.
    ↪format(theta[0][0],theta[1][0]))
print('Final cost/MSE:  {:0.3f}'.format(cost_history[-1]))
```


```
Theta0:          0.730,
Theta1:          0.482
Final cost/MSE:  418.472
```

1.2.1 Convergence graphs

```
[ ]: fig,ax = plt.subplots(figsize=(12,8))

ax.set_ylabel(r'J($\theta$)')
ax.set_xlabel('Iterations')
_=ax.plot(range(iterations),cost_history,'b.',label=r'cost: $J_\theta$')
ax.legend()
```

```
[ ]: <matplotlib.legend.Legend at 0x7fd5b3862908>
```




```
[ ]: theta_history[:,1].shape
```

```
[ ]: (1000,)
```

```
[ ]: fig,ax = plt.subplots(figsize=(12,8))
```

```
ax.set_ylabel('J(Theta)')
ax.set_xlabel('Iterations')
_=ax.plot(range(iterations),theta_history[:,0],'g.',label=r'$\theta_0$')
_=ax.plot(range(iterations),theta_history[:,1],'r.',label=r'$\theta_1$')
ax.legend()
```

```
[ ]: <matplotlib.legend.Legend at 0x7fd5a921b198>
```




1.2.2 Predicting GRE scores

```
[ ]: X2 = np.array([df1.values[340:,0]]).T
y2 = np.array([df1.values[340:,1]]).T
mean_X2 = np.sum(X2)/len(X2)
X2_fs = (X2-mean_X)/range_X
X2_b = np.c_[np.ones(len(X2_fs)),X2_fs]
y_predict = []
def validation_ds(X,y,theta):
    predict = np.dot(X,theta)
    global y_predict
    y_predict = predict
    diff = (predict - y)
    plt.scatter(np.array(range(len(X)))+340,diff)
    plt.xlabel(r'$x_i$')
    plt.ylabel('prediction error')
    plt.show()

validation_ds(X2_b,y2,theta)
print("GRE Score : ",340)
print("Chance Of Admission : ",y_predict[340-340], " Error :_
    ↳",y_predict[0]-y2[0])
print("GRE Score : ",390)
```

```
print("Chance Of Admission : ",y_predict[390-340], " Error :␣
→",y_predict[50]-y2[50])
```



```
GRE Score : 340
Chance Of Admission : [0.67976141] Error : [-0.07023859]
GRE Score : 390
Chance Of Admission : [0.6990253] Error : [0.0590253]
```

2 Linear Regression With Multiple Variable

```
[ ]: df1 = pd.DataFrame(df, columns=[df.columns.values[1]]+[df.columns.
→values[6]]+[df.columns.values[8]])
df1
```

```
[ ]:
GRE Score  CGPA  Chance of Admit
0         337   9.65             0.92
1         324   8.87             0.76
2         316   8.00             0.72
3         322   8.67             0.80
4         314   8.21             0.65
...        ...   ...             ...
395        324   9.04             0.82
396        325   9.11             0.84
397        330   9.45             0.91
```

398	312	8.78	0.67
399	333	9.66	0.95


[400 rows x 3 columns]

2.1 Plotting of Data:

```
[ ]: X1 = np.array([df1.values[:340,0]]).T
mean_X1 = np.sum(X1)/len(X1)
range_X1 = max(X1)-min(X1)
X1 = (X1-mean_X1)/range_X1

X2 = np.array([df1.values[:340,1]]).T
mean_X2 = np.sum(X2)/len(X2)
range_X2 = max(X2)-min(X2)
X2 = (X2-mean_X2)/range_X2

y = np.array([df1.values[:340,2]]).T
plt.scatter(X1,y)
plt.scatter(X2,y)
plt.xlabel(df1.columns.values[0])
plt.ylabel(df1.columns.values[1])
plt.show()
print(X.shape, y.shape)
```



(340, 1) (340, 1)

2.2 Gradient Descent:

```
[ ]: def cal_cost(theta,X,y):
    m = len(y)
    predictions = X.dot(theta)
    cost = (1/2*m) * np.sum(np.square(predictions-y))
    return cost

[ ]: def gradient_descent(X,y,theta,learning_rate=0.01,iterations=100):
    m = len(y)
    cost_history = np.zeros(iterations)
    theta_history = np.zeros((iterations,3))
    for it in range(iterations):
        prediction = np.dot(X,theta)
        # print(prediction.shape, y.shape)
        # print(np.array(prediction-y).shape)
        # print( X.T.dot((prediction - y)))
        theta = theta -(1/m)*learning_rate*( X.T.dot((prediction - y)))
        theta_history[it,:] =theta.T
        cost_history[it] = cal_cost(theta,X,y)
    return theta, cost_history, theta_history

[ ]: # theta = np.random.randn(2,1)
theta = np.array([[0],[0],[0]])
learning_rate = 0.1
iterations = 1000
X_b = np.c_[np.ones(len(X)),X1,X2]
# print(theta.shape, X_b[0], y[0])
# print("sas",X_b.shape)
theta,cost_history,theta_history = \
    ↪gradient_descent(X_b,y,theta,learning_rate,iterations)

print('Theta0:          {:.3f},\nTheta1:          {:.3f},\nTheta2:          {:.
    ↪0.3f}'.format(theta[0][0],theta[1][0],theta[2][0]))
print('Final cost/MSE:  {:.3f}'.format(cost_history[-1]))
```


Theta0: 0.730,
Theta1: 0.238,
Theta1: 0.383
Final cost/MSE: 270.959

2.2.1 Convergence graphs

```
[ ]: fig,ax = plt.subplots(figsize=(12,8))

ax.set_ylabel(r' $J(\theta)$ ')
ax.set_xlabel('Iterations')
_=ax.plot(range(iterations),cost_history,'b.',label=r'cost:  $J_\theta$ ')
ax.legend()
```

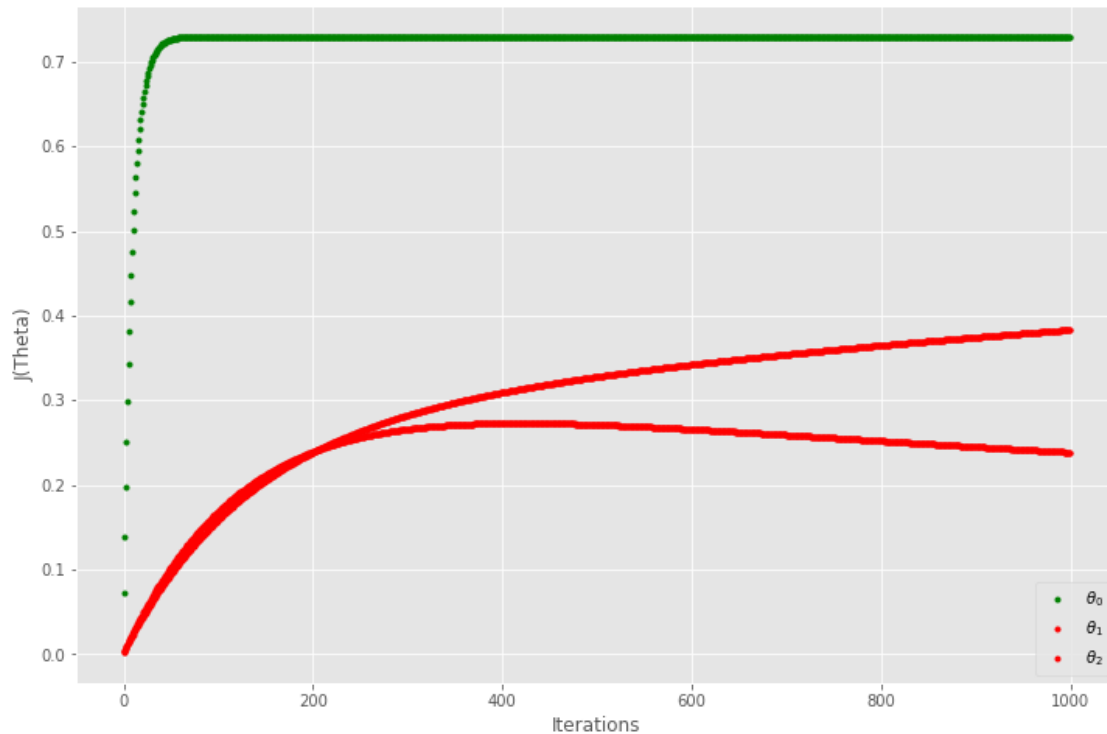
[]: <matplotlib.legend.Legend at 0x7fd5a9131ef0>



```
[ ]: fig,ax = plt.subplots(figsize=(12,8))

ax.set_ylabel('J(Theta)')
ax.set_xlabel('Iterations')
_=ax.plot(range(iterations),theta_history[:,0],'g.',label=r' $\theta_0$ ')
_=ax.plot(range(iterations),theta_history[:,1],'r.',label=r' $\theta_1$ ')
_=ax.plot(range(iterations),theta_history[:,2],'r.',label=r' $\theta_2$ ')
ax.legend()
```

[]: <matplotlib.legend.Legend at 0x7fd5a92b1630>



```
[ ]: X1_p = np.array([df1.values[340:,0]]).T
X1_p = (X1_p - mean_X1)/range_X1

X2_p = np.array([df1.values[340:,1]]).T
X2_p = (X2_p - mean_X2)/range_X2

y_p = np.array([df1.values[340:,2]]).T

X_b = np.c_[np.ones(len(X1_p)),X1_p,X2_p]
print(X_b.shape)
y_predict_2 = []
def validation_ds(X,y,theta):
    predict = np.dot(X,theta)
    global y_predict_2
    y_predict_2 = predict
    diff = (predict - y)
    plt.scatter(np.array(range(len(X)))+340,diff)
    plt.xlabel(r'$x_i$')
    plt.ylabel('prediction error')
    plt.show()


validation_ds(X_b,y_p,theta)
print("GRE Score : ",340)
```

```

print("Chance Of Admission : ",y_predict_2[340-340], " Error :␣
→",y_predict_2[0]-y_p[0])
print("GRE Score : ",390)
print("Chance Of Admission : ",y_predict_2[390-340], " Error :␣
→",y_predict_2[50]-y_p[50])

```

(60, 3)



```

GRE Score : 340
Chance Of Admission : [0.68387467] Error : [-0.06612533]
GRE Score : 390
Chance Of Admission : [0.66639327] Error : [0.02639327]

```

```

[ ]: 
[ ]: 

```

3 Logistic Regression

3.1 Loading Data

```
[ ]: data3=pd.read_csv('/content/drive/My Drive/Deep Learning/Lab1/heart.csv')
data3.head()
```

```
[ ]:   age  sex  cp  trestbps  chol  fbs  ...  exang  oldpeak  slope  ca  thal
target
0    63    1   3      145   233    1  ...    0      2.3      0   0    1
1
1    37    1   2      130   250    0  ...    0      3.5      0   0    2
1
2    41    0   1      130   204    0  ...    0      1.4      2   0    2
1
3    56    1   1      120   236    0  ...    0      0.8      2   0    2
1
4    57    0   0      120   354    0  ...    1      0.6      2   0    2
1

[5 rows x 14 columns]
```

3.2 Data Subset

```
[ ]: print(data3[['age', 'thalach', 'target']])
```

```
   age  thalach  target
0    63     150       1
1    37     187       1
2    41     172       1
3    56     178       1
4    57     163       1
..   ...     ...     ...
298  57     123       0
299  45     132       0
300  68     141       0
301  57     115       0
302  57     174       0
```

[303 rows x 3 columns]

3.3 Data / Feature Normalization

```
[ ]: X=data3[['age', 'thalach']].values
y=(data3.iloc[:, -1].values).reshape(-1,1)

X=(X-np.mean(X,0))/np.std(X,0)
```


```
print(X[:10,:])
```

```
[[ 0.9521966  0.01544279]
 [-1.91531289  1.63347147]
 [-1.47415758  0.97751389]
 [ 0.18017482  1.23989692]
 [ 0.29046364  0.58393935]
 [ 0.29046364 -0.07201822]
 [ 0.18017482  0.1466343 ]
 [-1.1432911  1.0212444 ]
 [-0.26098049  0.54020884]
 [ 0.29046364  1.0649749 ]]
```

3.4 Data Visualization

```
[ ]: #plt.figure(figsize=(12,8))
sns.scatterplot(data3['age'],data3['thalach'],hue=data3['target'])
plt.xlabel('Age')
plt.ylabel('Thalach')
plt.title('Age vs Thalach')
```

```
[ ]: Text(0.5, 1.0, 'Age vs Thalach')
```



3.5 Train - Test Data Split

```
[ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
→random_state=42)

X_train=np.append(np.ones([len(X_train),1]),X_train,axis = 1)
X_test=np.append(np.ones([len(X_test),1]),X_test,axis = 1)
```

3.6 Gradient Descent , Cost Function and Predict Function

```
[ ]: def logreg_cost(theta,x,y):
    m,n = x.shape
    cost = -(1 / m) * np.sum((y * np.log((1/(1+np.exp(-x.dot(theta)))))) + ((1
→y) * np.log(1-(1/(1+np.exp(-x.dot(theta)))))))
    return cost

def Gradient_Descent3(iterations,alpha,X_train,y_train):
    n,f = X_train.shape # of Rows and Features
    theta = np.zeros((f,1))
    history3=[]
    for i in range(iterations):
        grad=(1 / n) * np.dot(X_train.T, (1/(1+np.exp(-(X_train.dot(theta))))
→y_train)

        cost=logreg_cost(theta,X_train,y_train)

        theta = theta - (alpha*(grad)/len(X_train))

        history3.append([cost]+list(theta.flatten()))

    return history3

def predict(data,theta):
    y_hat=1/(1+np.exp(-(np.dot(data,theta))))

    y_hat=[1 if i > 0.5 else 0 for i in y_hat]

    return y_hat
```

3.7 Epochs and Optimal Theta Values

```
[ ]: # Parameters
iterations = 300
alpha = 0.01
```

```
history3= Gradient_Descent3(iterations,alpha,X_train,y_train)
```

```
theta=history3[-1][1:]  
print('Optimal Cost')  
print(history3[-1][0])  
print('Optimal Values for theta0,theta1 and theta2')  
print(theta)
```

Optimal Cost

0.6923773926317269


Optimal Values for theta0,theta1 and theta2

[0.0005995753564405086, -0.001594773158020073, 0.002833153476409542]

3.8 Convergence Graph

```
[ ]: plt.plot(np.linspace(1,300,300),np.array(history3)[: ,0],color='r')  
plt.xlabel('Iterations')  
plt.legend(['error','m1','m2'])
```


```
[ ]: <matplotlib.legend.Legend at 0x7fd59d089cc0>
```



```
[ ]: plt.plot(np.linspace(1,300,300),np.array(history3)[: ,1],color='g',label='m1')  
plt.plot(np.linspace(1,300,300),np.array(history3)[: ,2],color='b',label='m2')
```

```
plt.plot(np.linspace(1,300,300),np.array(history3)[: ,3],color='r',label='m2')
plt.xlabel('Iterations')
plt.legend(['m1','m2','C'])
```

[]: <matplotlib.legend.Legend at 0x7fd59d001160>



3.9 Model Performance on Training and Testing Data

```
[ ]: from sklearn.metrics import accuracy_score

y_pred1 = predict(X_train,theta)
score1=accuracy_score(y_train,y_pred1)
print('Training Accuracy :',score1)

y_pred2 = predict(X_test,theta)
score2=accuracy_score(y_test,y_pred2)
print('Testing Accuracy :',score2)
```

Training Accuracy : 0.6745283018867925
Testing Accuracy : 0.7142857142857143

3.10 Comparison with in-built Logistic Regression function

```
[ ]: from sklearn.linear_model import LogisticRegression

classifier = LogisticRegression(random_state = 0)
classifier.fit(X_train, y_train)

y_pred = classifier.predict(X_train)
print('Training Accuracy :', accuracy_score(y_train, y_pred))

y_pred = classifier.predict(X_test)
print('Testing Accuracy :', accuracy_score(y_test, y_pred))
```

Training Accuracy : 0.6933962264150944

Testing Accuracy : 0.7032967032967034

/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py:760:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
y = column_or_1d(y, warn=True)

```
[ ]:
```